

Tuesday 25/11

3pm - 5:30pm

Oracle-Efficient Combinatorial Semi-Bandits	Junghun Kim
OSKAR: Omnimodal Self-supervised Knowledge Abstraction and Representation	Kaouther Messaoud
A solvable model of learning generative diffusion: theory and insights	Hugo Cui
A Theoretical Framework for Grokking: Interpolation followed by Riemannian Norm Minimisation	Scott Pesme, Etienne Boursier
Algorithm- and Data-Dependent Generalization Bounds for Diffusion Models	Benjamin Dupuis, Dario Sharidian, Alain Oliviero-Durmus
AlligatOR: Pre-Training through Covisibility Segmentation for Relative Camera Pose Regression	Thibaut Loiseau
AlphaBeta is not as good as you think: a new probabilistic model to better analyze deterministic game-solving algorithms	Raphael Boige, Bruno Scherrer
Asymptotics of SGD in Sequence-Single Index Models and Single-Layer Attention Networks	Bruno Loureiro
Attention-based clustering	Rodrigo Maulen Soto
Bayes optimal learning of attention-indexed models	Fabrizio Boncoraglio
Bivariate Matrix-valued Linear Regression (BMLR): Finite-sample performance under Identifiability and Sparsity Assumptions	Nayel Bettache
Boosting Generative Image Modeling via Joint Image-Feature Synthesis	Theodoros Kouzelis
Bridging Arbitrary and Tree Metrics via Differentiable Gromov Hyperbolicity	Pierre Houdry, Titouan Vayer
Comparing Uniform Price and Discriminatory Multi-Unit Auctions through Regret Minimization	Marius Potfer
Convolution Goes Higher-Order: A Biologically Inspired Mechanism Empowers Image Classification	Simone Azeglio
EquiTaNPN: A Target-Permutation Equivariant Prior Fitted Network	Michael Arbel
Instance-Dependent Regret Bounds for Nonstochastic Linear Partial Monitoring	Khaled Eldowa
Large Stepsizes Accelerate Gradient Descent for Regularized Logistic Regression	Pierre Marion
Learning (Approximately) Equivariant Networks via Constrained Optimization	Luiz Chamom
Learning Theory for Kernel Bilevel Optimization	Fares El Khoury
Learning with Equality Constraints	Luiz Chamom
Learning with Restricted Boltzmann Machines: Asymptotics of AMP and GD in High Dimensions	Yizhou Xu
Memorization in Graph Neural Networks	Adarsh Jamadandi
Minimax Adaptive Online Nonparametric Regression over Besov Spaces	Paul Lioutaud
Multi-Token Prediction Needs Registers	Anastasios Gerontopoulos
Non-Asymptotic Analysis Of Data Augmentation For Precision Matrix Estimation	Lucas Morisset, Alain Oliviero-Durmus
Non-stationary Lipschitz Bandits	Nicolas Nguyen
On the Closed-Form of Flow Matching: Generalization Does Not Arise from Target Stochasticity	Anne Gagneux
Optimal Best Arm Identification under Differential Privacy	Marc Jourdan
Optimal Spectral Transitions in High-Dimensional Multi-Index Models	Leonardo Defilippis
Robust Distributed Estimation: Extending Gossip Algorithms to Ranking and Trimmed Means	Anna Van Elst
Sample and Map from a Single Convex Potential using Conjugate Generative Moment Measures	Nina Vesseron
ShiQ: Bringing back Bellman to LLMs	Pierre Clavier
SplitFlow: Flow Decomposition for Inversion-Free Text-to-Image Editing	Gaspard Beaudouin
T-REGS: Minimum Spanning Tree Regularization for Self-Supervised Learning	Julie Mordacq, David Loiseaux
TabArena: A Living Benchmark for Machine Learning on Tabular Data	David Holzmüller
Taxonomy of reduction matrices for Graph Coarsening	Antonin Joly
The quest for the GRAph Level autoEncoder (GRALE)	Paul Krzakala, Florence Dalche, Rémi Flamary
Towards Fully FP8 GEMM LLM Training at Scale	Dhia Garbaya
Uniform Generalization Bounds on Data-Dependent Hypothesis Sets via PAC-Bayesian Theory on Random Sets	Paul Viallard
Unsupervised Learning for Optimal Transport plan prediction between unbalanced graphs	Sonia Mazelet, Rémi Flamary
Why Diffusion Models Don't Memorize: The Role of Implicit Dynamical Regularization in Training	Raphaël Urfin
Streaming Federated Learning with Markovian Data	Tan-Khiem Huynh
Backward Conformal Prediction	Etienne Gauthier
Convergence of the Gradient Flow for Shallow ReLU Networks on Weakly Interacting Data	Léo Dana

Wednesday 26/11

3pm - 5:30pm

[Valid Selection among Conformal Sets	Mahmoud Hegazy
Learning long range dependencies through time reversal symmetry breaking	Maxence Errnout
A Markov Decision Process for Variable Selection in Branch & Bound	Paul Strang
Continuous Simplicial Neural Networks	Aref Einizade
Assessing the quality of denoising diffusion models in Wasserstein distance: noisy score and optimal bounds	Vahan Arsenyan, Elen Vardanyan
CLIPPTA: Robust Contrastive Vision-Language Test-Time Adaptation	Marc Lafon
Connecting Jensen–Shannon and Kullback–Leibler Divergences: A New Bound for Representation Learning	Reuben Dorent
Curl Descent : Non-Gradient Learning Dynamics with Sign-Diverse Plasticity	Hugo Ninou
Decomposing stimulus-specific sensory neural information via diffusion models	Simone Azeglio, Stéeve Laquitaine
Decreasing Entropic Regularization Averaged Gradient for Semi-Discrete Optimal Transport	Ferdinand Genans
DINO-Foresight: Looking into the Future with DINO	Efstathios Karypidis
Enhancing Optimizer Stability: Momentum Adaptation of The NGN Step-size	Rustem Islamov
ENMA: Tokenwise Autoregression for Continuous Neural PDE Operators	Lise Le Boudec
Fine-Tuning Discrete Diffusion Models with Policy Gradient Methods	Oussama Zekri
Learning Latent Variable Models via Jarzynski-adjusted Langevin Algorithm	Davide Carbone
Learning single index models via harmonic analysis	Hugo Koubbi
Local Curvature Descent: Squeezing More Curvature out of Standard and Polyak Gradient Descent	Simone Maria Giancola
MAP Estimation with Denoisers: Convergence Rates and Guarantees	Scott Pesme
Multi-Modal Attention Framework for Underwater Bioacoustic Denoising and Recognition	Amine Razig
On Logic-based Self-Explainable Graph Neural Networks	Alessio Ragno, Marc Plantet, Céline Robardet
Optimal Regret of Bandits under Differential Privacy	Achraf Azize
Enhancing Graph Classification Robustness with Singular Pooling.	Yassine Abbahaddou

A Geometry-Aware Metric for Mode Collapse in Multivariate Time Series Generative Models.	Yassine Abbahaddou
Relaxing partition admissibility in Cluster-DAGs: a causal calculus with arbitrary variable clustering	Clément Yvernes
Riemannian Flow Matching for Brain Connectivity Matrices via Pullback Geometry	Antoine Collas, Ce Ju
RNNs perform task computations by warping neural representations	Arthur Pellegrino
Scalable Feature Learning on Huge Knowledge Graphs for Downstream Machine Learning	Félix Lefebvre, Gaël Varoquaux
Searching Latent Program Spaces	Clément Bonnet
Sparse Autoencoders Learn Monosemantic Features in Vision-Language Models	Quentin Bouliot
Stochastic Optimization in Semi-Discrete Optimal Transport: Convergence Analysis and Minimax Rate	Ferdinand Genans
Template-Guided 3D Molecular Pose Generation via Flow Matching and Differentiable Optimization	Noémie Bergues, Arthur Carre
Tile-level Histopathology image Understanding benchmark	Sofiène Boutaj, Pierre Marza, Stergios Christodoulidis, Maria Vakalopoulou
Toward Interpretable Evaluation Measures for Time Series Segmentation	Félix Chavelli
Tree Ensemble Explainability through the Hoeffding Functional Decomposition and TreeHFD Algorithm	Clément Bénard
Variational Inference with mixture of isotropic Gaussian	Marguerite Petit-Talamon
Wasserstein convergence of critically damped Langevin diffusions	Antonio Ocello, Stanislas Strasman
Thresholds for sensitive optimality and Blackwell optimality in stochastic games	Julien Grand-Ciément
Continuous Simplicial Neural Networks	Jhony H. Giraldo
Solar-GECO: Perovskite Solar Cell Property Prediction with Geometric-Aware Co-Attention	guangjie li
Identifiability of Deep Polynomial Neural Networks	Jayneel Parekh
Learning to Steer: Input-dependent Steering for Multimodal LLMs	Konstantin Usevich